

Leland Wendel

CS-260

Final Project (Graph)

I used the files found [here](https://github.com/Joseph-I-Jess/cs260_winter_2024/tree/main/module_final) as a project template.

([https://github.com/Joseph-I-Jess/cs260\\_winter\\_2024/tree/main/module\\_final](https://github.com/Joseph-I-Jess/cs260_winter_2024/tree/main/module_final))

## 1. A function to add a new vertex to the graph

```

18  // add node
19  // takes a single parameter "new_node"
20  void add_node(GraphNode *new_node)
21  {
22      nodes.push_back(new_node); // adds element new_node to the 'nodes' vector
23  }
```

The `add_node` function takes a single parameter “`new_node`”, which is a pointer to a `GraphNode` object. “`new_node`” represents the node that will be added to the graph. At line 22, the `push_back()` function from the vector class adds the element to the end of the vector “`nodes`”. This function essentially adds a new node to the graph.

### Testing/Evaluation:

The `GraphNode` struct is defined as:

```

9   struct GraphNode
10  {
11      std::string name;
12      std::vector<Edge *> neighbors;
13  };
```

First, I need to create `GraphNode` objects:

```

14  // create nodes
15  GraphNode A{"Corvallis"};
16  GraphNode B{"Salem"};
17  GraphNode C{"Eugene"};
18  GraphNode D{"Portland"};
19  GraphNode E{"Mt Hood"};
20  GraphNode F{"Bend"};
```

The `GraphNode` objects *A – F* are initialized with the corresponding names: Corvallis, Salem, Eugene, Portland, Mt Hood, and Bend . These objects are defined locally in the main function.

Next, I need to add nodes A-F to the graph:

```
25     g.add_node(&A);
26     g.add_node(&B);
27     g.add_node(&C);
28     g.add_node(&D);
29     g.add_node(&E);
30     g.add_node(&F);
```

Here, 'g' is an instance of the `Graph` class. The `add_node` function takes a pointer to `GraphNode` and appends it to the 'nodes' vector, adding the node to the graph.

Next, I need to print the graph's current nodes to verify that everything is working:

```
119     // print current nodes in the graph
120     void print_nodes()
121     {
122         cout << "\nCurrent nodes in the graph:" << endl;
123         for (GraphNode *node : nodes)
124         {
125             cout << node->name << endl;
126         }
127     }
```

The 'print\_nodes()' function iterates through the nodes vector and prints the name of each node.

Finally, call the 'print\_nodes()' function in main to print the current nodes of the graph:

```
Current nodes in the graph:
Corvallis
Salem
Eugene
Portland
Mt Hood
Bend
```

All nodes have been successfully added to the graph!  
This matches my result from the design phase.

## 2. A function to add a new edge between two vertices of the graph

```

25  // add edge
26  // takes a single parameter "new_edge" (pointer to edge object)
27  void add_edge(Edge *new_edge)
28  {
29      // access source member of the new_edge object
30      // add new_edge pointer to the neighbors vector of source
31      new_edge->source->neighbors.push_back(new_edge);
32  }

```

The `add_edge` function takes one parameter “`new_edge`”, which is a pointer to an `Edge` object. “`new_edge`” represents the edge that will be added to the graph. At line 31, the source member of the `new_edge` object is accessed and the `push_back()` function from the vector class adds the element to the end of the vector “`neighbors`” of the source node. This function essentially records the connection from the source node to the destination node as an edge.

### Testing/Evaluation:

The `Edge` struct is defined as:

```

7  struct Edge {
8      int weight;
9      GraphNode *source;
10     GraphNode *destination;
11 };

```

First, I need to create an `Edge` object:

```

34  // create edges
35  Edge AB{36, &A, &B};
36  Edge BA{36, &B, &A};
37  Edge AC{47, &A, &C};
38  Edge CA{47, &C, &A};
39  Edge AF{128, &A, &F};
40  Edge FA{128, &F, &A};

```

The `Edge` objects are created and named `AB`, `BA`, `AC`, `CA`, `AF`, & `FA`. These objects represent a connection between two nodes on the graph and each have a defined weight (distance in this case). The elements within the brackets are initialized to the `Edge` object, and they correspond to members of the `Edge` struct (shown above). On line 35, for example, ‘36’ initializes the weight member of the struct, ‘&A’ initializes the source member of the edge struct, and ‘&B’ initializes the destination member of the edge struct.

```

59     // hook up the edges into the nodes
60     g.add_edge(&AB);
61     g.add_edge(&BA);
62     g.add_edge(&AC);
63     g.add_edge(&CA);
64     g.add_edge(&AF);
65     g.add_edge(&FA);

```

Here, 'g' is an instance of the Graph class. The add\_edge function adds an edge to the graph by updating the source node's list of neighbors. It takes a pointer to an Edge object and appends that Edge to the neighbors vector.

Next, I need to verify that the edges have been added:

```

81     // source, destination complication might come up here!
82     cout << "A.neighbors[2]->destination->name: " << A.neighbors[2]->destination->name << endl;
83     cout << "A.neighbors[1]->destination->name: " << A.neighbors[1]->destination->name << endl;
84     cout << "A.neighbors[0]->destination->name: " << A.neighbors[0]->destination->name << endl;

```

A.neighbors[2]->destination->name: Bend  
A.neighbors[1]->destination->name: Eugene  
A.neighbors[0]->destination->name: Salem

The edges have been successfully added to the graph! (at least for this example)  
This result matches the expected result from my design

### 3. A function for a shortest path algorithm

*\*Broken up into smaller sections:*

```

34     // shortest path (Dijkstra's algorithm)
35     // returns a vector of 'pair' objects where each pair consists of a GraphNode and an int
36     // dijkstra() takes one argument: a pointer to the starting GraphNode
37     vector<pair<GraphNode *, int>> dijkstra(GraphNode *start)
38     {
39         // create a hash table (distances) that maps each GraphNode to an int
40         // store the shortest known distance to each node from start node
41         unordered_map<GraphNode *, int> distances;
42
43         // initializes 'distances' for each node
44         for (GraphNode *node : nodes)
45         {
46             // distances to each node is infinity (INT_MAX)
47             distances[node] = INT_MAX;
48         }
49         // start distance should be zero (distance to itself)
50         distances[start] = 0;
51
52         // lambda function 'compare' to compare pairs (left vs right)
53         auto compare = [](pair<int, GraphNode *> left, pair<int, GraphNode *> right)
54         { return left.first > right.first; }; // if 1st ele of left pair > 1st ele of right pair - return true
55         // priority queue (min-heap) to store pairs (GraphNode & int)
56         // ordered using 'compare'
57         priority_queue<pair<int, GraphNode *>, vector<pair<int, GraphNode *>>, decltype(compare)> pq(compare);
58         // push start node on top of priority queue (PQ) because distance is zero
59         pq.push({0, start});

```

This function returns a vector of 'pair' objects, consisting of a GraphNode (GN) pointer and an integer. The pair represents the shortest distance from the start node to each node in the graph. The Dijkstra function takes one argument 'start', which is a pointer to the starting GraphNode. At line 41, a hash table 'distances' is created to map each GN pointer to an int (distance). The shortest known distance from the start node to each node will be stored here. At line 44, a for loop iterates over distances and sets each node's distance to infinity (or just a really large number). At line 50, the starting node's distance is set to zero (or distance to itself). At lines 53 and 54, a 'compare' function is declared to compare pairs – returns true if the first element of the left pair is greater than the first element of the right pair (constructor for min-heap). At line 57, the priority queue is declared to store the pairs of 'int' and GN and it is ordered using the compare function. At line 59, the start node is pushed on top of the PQ with a distance of zero.

```

61 // loop until PQ is empty
62 while (!pq.empty())
63 {
64     // extract node with smallest distance from PQ
65     auto [current_distance, current_node] = pq.top(); // structured binding unpacks pair returned by 'pq.top()'
66     // pop off PQ
67     pq.pop();
68
69     // if current distance > known/recorded distance:
70     if (current_distance > distances[current_node])
71     {
72         // skip this node
73         continue;
74     }
75     // iterate through each neighbor (edge connected) of current_node
76     for (Edge *edge : current_node->neighbors)
77     {
78         GraphNode *neighbor = edge->destination;
79         // calculate distance to neighboring node for each edge
80         int distance = current_distance + edge->weight;
81
82         // if new distance < known/recorded distance
83         if (distance < distances[neighbor])
84         {
85             // update distance
86             distances[neighbor] = distance;
87             // push new distance and node onto PQ
88             pq.push({distance, neighbor});
89         }
90     }
91 }

```

At line 62, a while loop is set to run until the PQ is empty. At line 65, the node with the smallest distance is extracted from the PQ using a structured binding that unpacks the pair returned by 'pq.top()'. At line 70, the node is skipped if its current distance is greater than the distance already recorded. At line 75, a for loop iterates through neighbors (connected by an edge) of the current node and calculates the distance to each neighboring node. At line 82, if the calculated distance to the neighbor is less than the known distance, the distance is updated, and the node and new distance is pushed onto the PQ.



```

92     // declares a vector 'result' to store pair of (GraphNode and int)
93     // int represents shortest distance from start node to GraphNode pointer
94     vector<pair<GraphNode *, int>> result;
95     // iterate over distances hash table
96     for (auto &[node, distance] : distances) // unpack each pair from distances into 'node' and 'distance'
97     {
98         // push a new pair to the end of the result vector
99         result.push_back({node, distance});
100    }
101    return result;
102 }

```

*After exiting the while loop (PQ is empty), a result vector is constructed from the distances map. The result vector contains pairs of each node and its corresponding shortest distance from the start node.*

### Testing/Evaluation:

*I need to test the algorithm and compare the results to my design's results:*

```

86     // Test Dijkstra's algorithm from Corvallis
87     cout << "\nDijkstra's shortest paths from Corvallis:\n";
88     vector<pair<GraphNode *, int>> distances = g.dijkstra(&A);
89     for (auto &[node, distance] : distances)
90     {
91         cout << node->name << ": " << distance << endl;
92     }
93
94     // Test Dijkstra's algorithm from Eugene
95     cout << "\nDijkstra's shortest paths from Bend:\n";
96     distances = g.dijkstra(&F);
97     for (auto &[node, distance] : distances)
98     {
99         cout << node->name << ": " << distance << endl;
100    }

```

*The Dijkstra function returns a vector of pairs (pointer to a GN and an int). At line 88, the function is called on graph 'g', starting from the node pointed to by '&A' (Corvallis). In the for loop, each pair is unpacked into 'node' and 'distance' and then each pair is printed to the console. At line 96, this process is repeated starting from the node pointed to by '&F' (Bend). The results are shown below:*

```
Dijkstra's shortest paths from Corvallis:
Bend: 128
Mt Hood: 159
Portland: 82
Eugene: 47
Salem: 36
Corvallis: 0
```

```
Dijkstra's shortest paths from Eugene:
Bend: 129
Mt Hood: 189
Portland: 112
Eugene: 0
Salem: 66
Corvallis: 47
```

*The results from the algorithm match the results from my design!*

## 4. A function for a minimum spanning tree algorithm

*Broken up into smaller sections:*

```
104 // minimum spanning tree (Prim's algorithm)
105 // returns a vector of pointers to Edge objects
106 vector<Edge*> prim()
107 {
108     // check if the graph has nodes
109     if (nodes.empty())
110         // if true, return empty vector
111         return {};
112     // inMST is a hash table to track if nodes are to be included in MST
113     unordered_map<GraphNode*, bool> inMST; // track if included, True/False
114     // vector to store edges that are part of the MST
115     vector<Edge*> mst;
116
117     // lambda function 'compare' to prioritize edges with smaller distances
118     auto compare = [](pair<int, Edge*> left, pair<int, Edge*> right)
119     { return left.first < right.first; };
120     // PQ to store pairs (pointers to edges and edge weights) ordered with compare function
121     priority_queue<pair<int, Edge*>, vector<pair<int, Edge*>>, decltype(compare)> pq(compare);
122
123     // helper function
124     // lambda function to add all edges associated with the node to PQ, if not in MST
125     auto add_edges = [&pq, &inMST](GraphNode* node)
126     {
127         // node included in MST = true
128         inMST[node] = true;
129         // for each Edge pointer in the neighbors vector of node
130         for (Edge* edge : node->neighbors)
131         {
132             // if destination node of the edge is not in MST
133             if (!inMST[edge->destination])
134             {
135                 // then add edge pair into PQ
136                 pq.push({edge->weight, edge});
137             }
138         }
139     };
140 }
```

The `prim` function returns a vector of pointers to `Edge` objects, which represent the edges in the minimum spanning tree (MST). At lines 109-111, I check if the graph is empty and return an empty vector if `nodes.empty()` is true. At lines 113 and 114, a hash table (`inMST`) is created to keep track of whether a node is included in the MST (true if it is, false if it isn't) and a vector (`MST`) is created to store edges that are a part of the MST. At line 125, the lambda function '`compare`' defines the comparison for the priority queue (PQ), prioritizing edges with smaller weights (shorter distances). The function takes two parameters '`left`' and '`right`', both of which are pairs of (`int`, `Edge *`). At line 119, the statement is true if the first element (weight) of the '`left`' pair is greater than the first element of the '`right`' pair. At line 121, a PQ is created to store the pairs (`weight`, `Edge *`) that have been ordered by weight using the '`compare`' function. At line 125, the lambda function '`add_edges`' adds all edges from the given node to the PQ if the destination node is not already in MST. This function marks the node as included in `inMST` and iterates over the node's neighbors to push other eligible edges into the PQ.

```

141         // Start with the first node
142         // add edges to PQ
143         add_edges(nodes[0]);
144
145         // loop until PQ is empty and MST contains 'n-1' edges
146         while (!pq.empty() && mst.size() < nodes.size() - 1)
147         {
148             // extracts top element of PQ (weight and pointer to Edge)
149             auto [weight, edge] = pq.top(); // structured binding
150             pq.pop(); // removes element
151
152             // if destination node is already in MST (true in the inMST map)
153             if (inMST[edge->destination])
154                 // skip current iteration
155                 continue;
156
157             // edge is added to the MST
158             mst.push_back(edge);
159             // mark destination node as included in MST
160             // add all its edges (not yet in MST) to the PQ
161             add_edges(edge->destination);
162         }
163
164         return mst;
165     }

```

At line 143, MST is initialized by calling '`add_edges`' for the first node of the graph, adding its edges to the PQ. The while loop at line 146 will run until the PQ is empty and the MST contains '`n-1`' edges (an MST for a connected graph contains '`n-1`' edges). At lines 149 and 150, the edge on top of the PQ is selected and removed from the PQ. At lines 153 and 155, if the destination node of the edge is already in the MST, the edge is skipped. At line 158, if the edge is not in MST, then edge is added to MST. Finally, at line



161, the 'add\_edges' function is called for the destination node to add its edges to the PQ.

### Testing/Evaluation:

```

102 // Test Prim's algorithm for MST
103 cout << "\nPrim's minimum spanning tree:\n";
104 vector<Edge *> mst = g.prim();
105 for (Edge *edge : mst)
106 {
107     cout << "(" << edge->source->name << " - " << edge->destination->name << ", " << edge->weight << ")\n";
108 }

```

The prim function returns a vector 'MST' containing edge pointers which represent the minimum spanning tree. The prim function is called on the current graph and the for loop iterates over each edge in MST. The details of each edge in the MST are then printed.

A second smaller graph was created to provide a second test:

```

110 // Create a second graph
111 Graph g2;
112 GraphNode G{"Brookings"};
113 GraphNode H{"Coos Bay"};
114 GraphNode I{"Roseburg"};
115 GraphNode J{"Medford"};
116
117 g2.add_node(&G);
118 g2.add_node(&H);
119 g2.add_node(&I);
120 g2.add_node(&J);
121
122 Edge GH{107, &G, &H};
123 Edge HG{107, &H, &G};
124 Edge GI{162, &G, &I};
125 Edge IG{161, &I, &G};
126 Edge HJ{165, &H, &J};
127 Edge JH{165, &J, &H};
128 Edge IJ{97, &I, &J};
129 Edge JI{97, &J, &I};
130
131 g2.add_edge(&GH);
132 g2.add_edge(&HG);
133 g2.add_edge(&GI);
134 g2.add_edge(&IG);
135 g2.add_edge(&HJ);
136 g2.add_edge(&JH);
137 g2.add_edge(&IJ);
138 g2.add_edge(&JI);
139
140 cout << "\nPrim's minimum spanning tree (second graph):\n";
141 mst = g2.prim();
142 for (Edge *edge : mst)
143 {
144     cout << "(" << edge->source->name << " - " << edge->destination->name << ", " << edge->weight << ")\n";
145 }

```

```

Prim's minimum spanning tree:
(Corvallis - Salem, 36)
(Salem - Portland, 46)
(Corvallis - Eugene, 47)
(Portland - Mt Hood, 77)
(Corvallis - Bend, 128)

Prim's minimum spanning tree (second graph):
(Brookings - Coos Bay, 107)
(Brookings - Roseburg, 162)
(Roseburg - Medford, 97)

```

*The final result matches my prediction from the design portion!*

## 5. Analyze the complexity of all of your graph behaviors

*The `add_node` and `add_edge` functions both have a complexity of  $O(1)$  if their vectors don't need resizing, in which case the complexity is  $O(n)$ .*

*In the Dijkstra function:*

- *The distance initialization has a complexity of  $O(n)$ , where 'n' is the number of nodes in the graph.*
- *The PQ initialization and push has a complexity of  $O(\log n)$*
- *Edge processing occurs once for each edge, so  $O(n)$  where  $n$  is the number of edges*
- *For each edge, the PQ operation is  $O(\log n)$ .*

*Overall, the algorithm has a complexity of  $O(E \log N)$ , where  $E$  is the number of edges and  $N$  is the number of nodes .*

*In the prim function:*

- *Empty graph check is  $O(1)$*
- *Hash table and MST vector initialization is  $O(1)$*
- *PQ initialization is  $O(1)$*
- *Add\_edges function is  $O(\log n)$ , where  $n$  is the number of edges*

*Overall, the complexity of the prim function is  $O(E \log N)$ , where  $E$  is the number of edges and  $N$  is the number of nodes.*